Hypothesis Testing in Data Science

A Review of Methods, Applications, and Emerging Practices

Sidney J Reynaud, Jr.

National University

Literature Review

July 19, 2026

Abstract

Hypothesis testing remains a foundational component of inferential statistics and modern data science. As data environments expand in scale and complexity, traditional null hypothesis significance testing (NHST) faces increasing scrutiny for p-value misuse, limited interpretability, and vulnerability to reproducibility failures. Contemporary practice incorporates resampling methods, false discovery rate (FDR) control, sequential testing, and Bayesian inference to address these limitations. This condensed review synthesizes classical foundations, critiques of NHST, and emerging methodologies relevant to large-scale data science workflows. As noted in my document, "Hypothesis testing occupies a foundational role in inferential statistics …" (Fisher, 1935; Neyman & Pearson, as cited in Wasserstein & Lazar, 2016).

Introduction

Hypothesis testing has shaped quantitative inquiry for more than a century, originating in early probability theory and being formalized by Fisher (1935), Neyman, and Pearson. NHST became the dominant inferential framework, relying on p-values to evaluate whether observed results are unlikely under a null hypothesis. However, widespread misinterpretation, particularly the treatment of p-values as the probability that the null hypothesis is true, has contributed to reproducibility failures across psychology, medicine, and the social sciences (Wasserstein & Lazar, 2016).

The replication crisis intensified scrutiny of NHST, prompting the American Statistical Association to issue guidance discouraging mechanical reliance on significance thresholds and encouraging broader inferential tools such as effect sizes, confidence intervals, and Bayesian methods. As the document states, “The crisis has prompted … formal guidance on the appropriate and limited role of p-values …” (Wasserstein & Lazar, 2016).

Classical Hypothesis Testing Foundations

Classical hypothesis testing begins with defining a null and alternative hypothesis, selecting a test statistic, and evaluating the probability of observing results as extreme as those obtained under the null (Fisher, 1935). Fisher’s early work established randomization and significance testing as core scientific tools, while Neyman and Pearson formalized Type I and Type II errors.

Cohen (1988) emphasized chronically underpowered studies and highlighted effect sizes as essential complements to significance tests. His work demonstrated that statistical significance alone provides limited insight without understanding the magnitude of an effect.

Modern data science exposes limitations of classical testing. Extremely large samples can make trivial effects statistically significant, while small or noisy datasets risk false negatives. As the document notes, “Very large samples … afford such statistical power that even trivially small effects attain significance …” (Cohen, 1988; Efron & Hastie, 2016).

Hypothesis Testing in Data Science Applications

Traditional Tests in Automated Pipelines

Tests such as t-tests, chi-square tests, and ANOVA remain widely used due to their efficiency and interpretability. They are embedded in automated workflows, model evaluation routines, and feature selection processes. However, their validity depends on assumptions of normality, independence, and homogeneity of variance that often fail in heterogeneous or observational datasets (Cohen, 1988; Efron & Hastie, 2016).

A/B Testing and Online Experimentation

A/B testing has become a cornerstone of digital experimentation. Kohavi et al. (2009) documented how large technology companies adapt classical tests to complex web environments characterized by skewed metrics, heterogeneous users, and network effects. Pre-experiment power analysis and minimum detectable effect calculations are now standard practice.

A major operational challenge is sequential monitoring (“peeking”), which inflates Type I error rates when practitioners stop experiments early based on interim significance. Johari et al. (2017) introduced always-valid inference and confidence sequences to address this issue, enabling continuous monitoring without compromising error control. As the document states, “A particular challenge … is the problem of continuous or sequential monitoring …” (Johari et al., 2017).

Resampling and Non-Parametric Approaches

Resampling methods provide robust alternatives when parametric assumptions fail. The bootstrap approximates sampling distributions by repeatedly resampling with replacement, a technique made feasible by modern computational power (Efron & Hastie, 2016). These methods are especially useful in machine learning contexts where analytical distributions are unavailable.

Permutation tests construct null distributions by shuffling labels under the assumption of exchangeability. Good (2005) demonstrated that permutation tests offer exact Type I error control and are well suited for small samples, heavy-tailed distributions, or unconventional test statistics. As the document notes, “Permutation tests … offer exact Type I error control under the null hypothesis …” (Good, 2005).

Multiple Testing and False Discovery Rate Control

Large-scale testing common in genomics, neuroimaging, and high-dimensional modeling creates substantial risk of false positives. Bonferroni correction controls the family-wise error rate but is overly conservative when thousands of tests are performed.

Benjamini and Hochberg (1995) introduced the false discovery rate (FDR), which controls the expected proportion of false positives among rejected hypotheses. Their procedure ranks p-values and identifies a cutoff that balances discovery and error control. As the document states, “This approach achieves a substantially more favorable power-to-error tradeoff …” (Benjamini & Hochberg, 1995).

Bayesian Alternatives

Bayesian hypothesis testing reframes inference by evaluating the probability of hypotheses given observed data. Bayes factors quantify relative evidence between competing models and can support the null hypothesis, something p-values cannot do (Held & Ott, 2018).

The ASA’s statement (Wasserstein & Lazar, 2016) encouraged broader inferential practices, including Bayesian methods, effect sizes, and confidence intervals. Held and Ott (2018) showed that p-values often exaggerate evidence against the null compared to Bayes factors, especially in small samples. As the document notes, “p-values tend to overstate the evidence against the null …” (Held & Ott, 2018).

Table 1

Summary of Hypothesis Tests Identified in the Literature Review.

Hypothesis Test	Application in Data Science	Assumptions	Key Strengths	Limitations	Sources
Two-Sample Proportion Test	Comparing conversion rates between Version A and Version B (A/B testing).	Independent samples; binary outcomes; adequate sample size.	Simple, interpretable; ideal for conversion analysis.	Sensitive to small sample sizes; cannot adjust for covariates.	Hypothesis_Tests.pdf; Kohavi et al. (2009)
Independent Samples t-Test	Comparing mean order totals or mean session duration across groups.	Normality; independence; equal/unequal variance depending on variant.	Fast, widely used; robust with Welch correction.	Sensitive to outliers; assumes continuous data.	Field (2018); Practical Application.pdf
Chi-Square Test of Independence	Testing association between categorical variables (e.g., Region × Purchase Completion).	Expected cell counts ≥ 5; independence.	Ideal for categorical comparisons; easy to compute.	Cannot measure strength of association without effect size (Cramér's V).	Agresti (2013); Practical Application.pdf
Permutation Test	Non-parametric comparison of means or proportions; useful when assumptions fail.	Exchangeability; representative sample.	Exact Type I error control; robust to non-normality.	Computationally intensive for large datasets.	Good (2005); Literature Review.pdf
Bootstrap Test	Estimating sampling distributions for complex ML metrics or non-analytic statistics.	Representative sample; independence.	Flexible; works with arbitrary statistics.	Can fail with biased or unrepresentative samples.	Efron & Hastie (2016)
False Discovery Rate (FDR)	Large-scale feature selection, genomics, ML model comparison.	Independent or positively dependent tests.	Controls proportion of false positives; more powerful than Bonferroni.	Not suitable when strict family-wise error control is required.	Benjamini & Hochberg (1995)
Sequential / Always-Valid Tests	Online A/B testing with continuous monitoring.	Martingale or confidence sequence assumptions.	Allows real-time decision-making without inflating Type I error.	More complex to implement; requires specialized statistical tools.	Johari et al. (2017)
Bayesian Hypothesis Testing	Model comparison, ML evaluation, decision-making under uncertainty.	Requires prior specification; independence.	Provides direct evidence for hypotheses; supports null hypothesis.	Sensitive to prior choice; computationally heavy.	Held & Ott (2018)

Note. This table summarizes the primary hypothesis tests identified in the literature review, including their applications in data science, underlying assumptions, strengths, limitations, and corresponding scholarly sources. Tests include classical parametric methods, non-parametric resampling approaches, sequential inference procedures, and Bayesian hypothesis testing frameworks.

Comparative Analysis

Assumptions

- Classical tests require distributional assumptions and independence (Fisher, 1935; Cohen, 1988).
- Resampling methods require representativeness and exchangeability but no parametric form (Efron & Hastie, 2016; Good, 2005).
- Bayesian methods require priors, which can introduce subjectivity but also incorporate domain knowledge (Held & Ott, 2018).

Interpretability

- P-values are widely used but frequently misunderstood (Wasserstein & Lazar, 2016).
- Bayes factors provide direct evidence for or against hypotheses (Held & Ott, 2018).
- FDR-adjusted q-values offer interpretable error rates in large-scale testing (Benjamini & Hochberg, 1995).

Scalability

- Classical tests scale efficiently (Fisher, 1935).
- Sequential tests support continuous monitoring (Johari et al., 2017).
- Bootstrap and permutation methods can be computationally heavy (Efron & Hastie, 2016; Good, 2005).
- Bayesian models vary widely in computational cost (Held & Ott, 2018).

Error Risks

- NHST inflates Type I error under multiple testing (Benjamini & Hochberg, 1995).
- Sequential peeking accelerates false positives (Johari et al., 2017).
- Bayesian methods risk prior sensitivity (Held & Ott, 2018).
- Resampling methods fail when samples are unrepresentative (Good, 2005).

Conclusion

Hypothesis testing in data science requires methodological flexibility. Classical NHST provides foundational structure but must be supplemented with effect sizes, confidence intervals, and modern alternatives. Resampling methods offer robustness, FDR procedures improve large-scale inference, and Bayesian approaches provide richer evidence frameworks. Sequential testing innovations support real-time experimentation, and methodological pluralism remains essential for reproducible scientific practice. As the document states, “Based on the synthesis of this literature, several best practices emerge …” (Wasserstein & Lazar, 2016; Held & Ott, 2018).

References

Benjamini, Y., & Hochberg, Y. (1995). *Controlling the false discovery rate: A practical and powerful approach to multiple testing*. *Journal of the Royal Statistical Society: Series B (Methodological), 57*(1), 289–300.

Cohen, J. (1988). *Statistical power analysis for the behavioral sciences* (2nd ed.). Lawrence Erlbaum Associates.

Efron, B., & Hastie, T. (2016). *Computer age statistical inference: Algorithms, evidence, and data science*. Cambridge University Press.

Fisher, R. A. (1935). *The design of experiments*. Oliver and Boyd.

Good, P. I. (2005). *Permutation, parametric, and bootstrap tests of hypotheses* (3rd ed.). Springer.

Held, L., & Ott, M. (2018). On p-values and Bayes factors. *Annual Review of Statistics and Its Application, 5*, 393–419.

Johari, R., Koomen, P., Pekelis, L., & Walsh, D. (2017). Peeking at A/B tests: Why it matters, and what to do about it. *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1517–1525.

Kohavi, R., Longbotham, R., Sommerfield, D., & Henne, R. L. (2009). Controlled experiments on the web: Survey and practical guide. *Data Mining and Knowledge Discovery, 18*(1), 140–181.

Wasserstein, R. L., & Lazar, N. A. (2016). The ASA's statement on p-values: Context, process, and purpose. *The American Statistician, 70*(2), 129–133.